

F1 Gate training loss. Val BCE is NOT comparable across arms: B and C carry deliberate subsampling
 A — no subsampling in their logits that Platt later removes, and BCE penalises exactly that bias. Compare AUC instead.
 best epoch 49 of 50

B — uniform 1/5
 best epoch 19 of 24

C — hard-neg $\propto 1/x^2$
 best epoch 5 of 10

