

F6 Before/after calibration on the val split. AUC-ROC and AUC-PR are IDENTICAL for raw and Platt-2 by construction ( $a > 0$  is monotone, and AUC sees only ranking); they move only under Platt-4, whose slope  $a(x) = a_0 + a_1 \cdot \log n_{\text{window}}$  is row-dependent. DR-ECE/ECE is how much the count-weighted ECE understates error where decisions happen — it is a diagnostic OF ECE, not a quality score: a HIGH ratio means the model is near-perfect on the 99.5% easy-negative bulk that dominates ECE, while a LOW ratio ( $\chi^2_{\lambda} \times 4$ , arm C row 2x) means the estimator is bad everywhere, so the bulk flatters it less.

